



Go For It

Building Advanced Systems with Go and Couchbase Server
7 October 2014

Marty Schoch






History

- Started small with Couchbase Labs
- Steady growth internally
- Now in production


Things we like

- First class concurrency support with clean code
- Custom structs easily map to JSON
- Out of the box support for HTTP/HTTPS





- Client SDK born in 2012
- Currently community supported
- Used for many projects internally
- Officially supported client soon





Intro to Couchbase SDK



○ Running Code in this Presentation



```
12 func handleError(err error) {  
13     if err != nil {  
14         panic(err)  
15     }  
16 }  
17  
18 func main() {  
19     err := doSomething()  
20     handleError(err)  
21     fmt.Printf("Success")  
22 }
```

Success
Program exited.

Run Kill Close



```
4 import (  
5     "fmt"  
6  
7     "github.com/couchbaselabs/go-couchbase"  
8 )  
9  
10 func main() {  
11     bucket, err := couchbase.GetBucket("http://localhost:8091/", "default", "demo")  
12     handleError(err)  
13     defer bucket.Close()  
14     fmt.Printf("Connected to Couchbase Bucket '%s'\n", bucket.Name)  
15 }
```

Connected to Couchbase Bucket 'demo'

Program exited.

Run Kill Close



```
8  type Event struct {
9      Type string `json:"type"`
10     Name string `json:"name"`
11     Likes int    `json:"likes"`
12 }
13
14 func NewEvent(name string) *Event {
15     return &Event{"event", name, 0}
16 }
17
18 func NewEventJSON(jsonbytes []byte) (event *Event) {
19     err := json.Unmarshal(jsonbytes, &event)
20     handleError(err)
21     return
22 }
23
24 func (e *Event) String() string {
25     return fmt.Sprintf("Event '%s', Likes: %d", e.Name, e.Likes)
26 }
```

- API supports working with any JSON serializable structure, or raw []byte
- Examples today will use the structure above



```
42  event := NewEvent("Couchbase Connect")
43  err = bucket.Set("cc2014", 0, event)
44  handleError(err)
45
46  event = NewEvent("GopherCon India")
47  err = bucket.Set("gci2015", 0, event)
48  handleError(err)
49
50  fmt.Printf("Saved Events\n")
```

Saved Events

Program exited.

Run Kill Close



```
42  var event Event
43  err = bucket.Get("cc2014", &event)
44  handleError(err)
45  fmt.Println(&event)
```

Event 'Couchbase Connect', Likes: 0

Program exited.

Run Kill Close



○ Mutation Ops



- Add/AddRaw
- Append
- Cas/CasRaw
- Delete
- Incr
- Set/SetRaw

These all have very similar semantics to the other SDKs.



- Upcoming Events
 - Couchbase Connect 2014 
 - GopherCon India 2015 

○ Liking Events - Wrong



```
38 func likeEvent(bucket *couchbase.Bucket, id string) {  
39     var event Event  
40     err := bucket.Get(id, &event)  
41     handleError(err)  
42     event.Likes++  
43     bucket.Set(id, 0, event)  
44 }
```

Concurrent Updates - Incorrect



```
54  var wg sync.WaitGroup
55  for i := 0; i < 100; i++ {
56      wg.Add(1)
57      go func() {
58          defer wg.Done()
59          likeEvent(bucket, "cc2014")
60      }()
61  }
62  wg.Wait()
63
64  var event Event
65  err = bucket.Get("cc2014", &event)
66  handleError(err)
67  fmt.Println(&event)
```

Event 'Couchbase Connect', Likes: 20

Program exited.

Run Kill Close

○ Liking Events - Right



```
38 func likeEvent(bucket *couchbase.Bucket, id string) {  
39     bucket.Update(id, 0, func(current []byte) ([]byte, error) {  
40         event := NewEventJSON(current)  
41         event.Likes++  
42         return json.Marshal(event)  
43     })  
44 }
```

○ Concurrent Updates - Safe using Update()



```
54     var wg sync.WaitGroup
55     for i := 0; i < 100; i++ {
56         wg.Add(1)
57         go func() {
58             defer wg.Done()
59             likeEvent(bucket, "cc2014")
60         }()
61     }
62     wg.Wait()
63
64     var event Event
65     err = bucket.Get("cc2014", &event)
66     handleError(err)
67     fmt.Println(&event)
```

Event 'Couchbase Connect', Likes: 120

Program exited.

Run Kill Close

○ Views - Top Events by Likes



```
1 function (doc, meta) {  
2     if(doc.type === 'event') {  
3         emit(doc.likes, null);  
4     }  
5 }
```

- Emit 1 row for every event
- Key is the number of likes
- No value, we just use this view to find Event IDs



```
21  args := map[string]interface{}{
22      "stale":      false,
23      "descending": true,
24  }
25
26  res, err := bucket.View("ddoc", "likes", args)
27  handleError(err)
28
29  for _, r := range res.Rows {
30      fmt.Printf("Key: %v - DocID: '%s'\n", r.Key, r.ID)
31  }
```

```
Key: 120 - DocID: 'cc2014'
Key: 0 - DocID: 'gci2015'
```

Program exited.

Run Kill Close

Behind the Curtains





```
import _ "github.com/couchbase/gomemcached/debug"
```

```
▼ "mc": {  
  ▼ "recv": {  
    ▼ "bytes": {  
      "ADD": 3789188,  
      "DELETE": 512073,  
      "GET": 255889947,  
      "GETQ": 245698,  
      "INCREMENT": 102850,  
      "SASL_AUTH": 935878,  
      "SASL_LIST_MECHS": 961172,  
      "SET": 14188813,  
      "STAT": 8245136,  
      "total": 284870755  
    },  
    "errs": {},  
    ▼ "ops": {  
      "ADD": 97504,  
      "DELETE": 21330,  
      "GET": 730624,  
      "GETQ": 620,  
      "INCREMENT": 3214,  
      "SASL_AUTH": 25294,  
      "SASL_LIST_MECHS": 25294,  
      "SET": 590796,  
      "STAT": 177768,  
      "total": 1672444  
    }  
  },  
  ▶ "tap": { ... }, // 3 items  
  ▼ "xmit": {
```

- Go stdlib hidden gem - <http://golang.org/pkg/expvar/>



- Operations on the Couchbase bucket ultimately need to talk to one of the Couchbase servers
- Bulk operations talk to multiple Couchbase servers
- Applications perform bucket operations on separate go routines, don't expect to be blocked by one another
- This is simulated by maintaining pools of connections to the underlying servers

○ Connection Pool Properties



- Return usable connection as fast as possible
- Creating connections is relatively expensive (as compared to reusing them)
- Don't create them unnecessarily
- Don't create too many of them
- The usual tuning operation here, too large a pool wastes resources, too small a pool means waiting for connections.



```
4  type connectionPool struct {
5      host      string
7      auth      AuthHandler
8      connections chan *memcached.Client
9      createsem  chan bool
10 }
11
12 func newConnectionPool(host string, ah AuthHandler, poolSize, poolOverflow int) *connectionPool {
13     return &connectionPool{
14         host:      host,
15         connections: make(chan *memcached.Client, poolSize),
16         createsem:  make(chan bool, poolSize+poolOverflow),
17         auth:      ah,
18     }
19 }
20 }
```

- Using a buffered channel of connections as a thread-safe pool
- Using a buffered channel of bools to track overflow connections

Quick Go Channel Review



```
8      // write to channel
9      channel <- val
10
11     // read from channel
12     val = <-channel
```

○ Connection Pool - Get Connection 1



```
4 func (cp *connectionPool) GetWithTimeout(d time.Duration) (rv *memcached.Client, err error)
5 {
6     // short-circuit available connections
7     select {
8     case rv, isopen := <-cp.connections:
9         if !isopen {
10             return nil, errClosedPool
11         }
12         return rv, nil
13     default:
14     }
```

- Select on the pool channel, if reading won't block, read and return connection
- If this would have blocked (no available connections in pool), proceed to next step

○ Connection Pool - Get Connection 2



```
18 // create a very short timer, 1ms
19 t := time.NewTimer(ConnPoolAvailWaitTime)
20 defer t.Stop()
21
22 select {
23 case rv, isopen := <-cp.connections:
24     // connection became available
25     if !isopen {
26         return nil, errClosedPool
27     }
28     return rv, nil
29 case <-t.C:
30     // waited 1ms
31 }
32 }
```

○ Connection Pool - Get Connection 3



```
46      t.Reset(d) // reuse original timer for full timeout
47      select {
48      case rv, isopen := <-cp.connections:
49
50          // keep trying to get connection from main pool
51          if !isopen {
52              return nil, errClosedPool
53          }
54          return rv, nil
55
56      case cp.createsem <- true:
57
58          // create a new connection
59          rv, err := cp.mkConn(cp.host, cp.auth)
60          if err != nil {
61              <-cp.createsem // buffer only allows poolSize + poolOverflow
62          }
63          return rv, err
64
65      case <-t.C:
66
67          // exceeded caller provided timeout
68          return nil, ErrTimeout
```

○ Connection Pool - Summary



- Somewhat dense block of Go code
- Worth your time to try to understand it
- This current version of the code was refined during performance benchmarks of `sync_gateway`
- See Dustin's blog

<http://dustin.sallings.org/2014/04/25/chan-pool.html>

Applications



○ cbugg - bug tracker on top of Couchbase



 **dustin**

bug-837 - consolio can survive cold reboots across any machine consolio
bug-828 - Set on authenticated buckets fail go-couchbase works-for-me
bug-827 - better logging of detailed events from gitmirror gitmirror
bug-801 - go-couchbase should support Append/Prepend() mutations go-couchbase
bug-741 - need an administrative view of all databases and gateways consolio
bug-677 - org-mode export of my bugs cbugg low-effort
bug-605 - automatic directory listings cbfs
bug-600 - Update the static assets whenever someone tries to load / cbgb ui
bug-578 - cbfs should use gorilla/mux cbfs debt
bug-576 - delete versions from phone home admin ui blocker current phone-home
bug-556 - couch rest API needs some auth work auth blocker cbgb
bug-466 - cbgb works as database for gamesim cbgb
bug-413 - replay append handling replay

- Typical CRUD operations, bugs, comments attachments
- Uses a large number of features in the SDK, but not a complex application



- Go HTTP server exposing REST API
- Also serves static resources HTML/CSS/JS/images
- End-user functionality through HTML5/AngularJS interface
- Bugs, Comments stored in Couchbase
- Searchable through Couchbase-Elasticsearch integration
- Attachments stored in cbfs



Ensure that the engineers building Couchbase rely on it being a high quality product.

○ cbugg - Deploying Views?



- 3-4 developers
- Important functionality built on top of views
- Each with local Couchbase, and shared production instance
- How do we propagate changes to design documents/views?
- Need to promote changes up to production, and back down to other developers

Version Controlled View Definitions



```
19 type viewMarker struct {
20     Version    int        `json:"version"`
21     Timestamp  time.Time   `json:"timestamp"`
22     Type       string     `json:"type"`
23 }
24
26 const ddockKey = "/@ddocVersion"
27 const ddocVersion = 1
28 const designDoc = `
29 {
30     "views": {
31         "likes": {
32             "map": "function (doc, meta) { if(doc.type === 'event') { emit(doc.likes, null);} }"
33         }
34     }
35 }`
```

- viewMarker tracks the latest deployed version
- we store viewMarker in *ddockKey*
- when we update *designDoc*, we bump the *ddocVersion*

○ Automatic View Definition Updating



```
50  marker := viewMarker{}
51  err := bucket.Get(ddocKey, &marker)
52  if err != nil && !gomemcached.IsNotFound(err) {
53      handleError(err)
54  }
55  if marker.Version < ddocVersion {
56      fmt.Printf("Installing new version of views (old version=%v)\n",
57          marker.Version)
58      doc := json.RawMessage([]byte(designDoc))
59      err = bucket.PutDDoc("ddoc", &doc)
60      handleError(err)
61      marker.Version = ddocVersion
62
63      marker.Timestamp = time.Now().UTC()
64      marker.Type = "ddocmarker"
65
66      bucket.Set(ddocKey, 0, &marker)
67  } else {
68      fmt.Printf("Version %v already installed\n", marker.Version)
69  }
```

Automatic View Definition Updating




```
26 const ddocKey = "/@ddocVersion"
27 const ddocVersion = 1
28 const designDoc = `
29 {
30   "views": {
31     "likes": {
32       "map": "function (doc, meta) { if(doc.type === 'event') { emit(doc.likes, null);} }"
33     }
34   }
35 }`
36
38
39 func main() {
40   bucket, err := couchbase.GetBucket("http
41   handleError(err)
42
43   updateDesignDocs(bucket)
44 }
```

Installing new version of views (old version=0)

Program exited.

Run Kill Close



- Server-side component integrating Couchbase Server and Couchbase Lite

○ Sync Gateway - How it Works



- Shared nothing architecture, need to scale Sync Gateway nodes just like Couchbase Server
- Sync Gateway maintains caches of data structures used for replication
- Relies on the Couchbase TAP protocol to be notified of changes
- These notifications invalidate/update cache

```

23  args := memcached.DefaultTapArguments()
24  feed, err := bucket.StartTapFeed(&args)
25  handleError(err)
26
27  go func() {
28      time.Sleep(1 * time.Second)
29      for i := 0; i < 5; i++ {
30          bucket.SetRaw(fmt.Sprintf("tap-%d", i), 0, []byte("x"))
31      }
32  }()
33
34  fmt.Printf("Listening to TAP:\n")
35  for op := range feed.C {
36      fmt.Printf("Received %s\n", op.String())
37      if len(op.Value) > 0 && len(op.Value) == 1 {
38          fmt.Printf("\tValue: %s\n", op.Value[0])
39      }
40  }

```

Listening to TAP:

Received <TapEvent Mutation, key="tap-4" (1 bytes
Value: x

Received <TapEvent Mutation, key="tap-3" (1 bytes
Value: x

Received <TapEvent Mutation, key="tap-2" (1 bytes
Value: x

Received <TapEvent Mutation, key="tap-0" (1 bytes
Value: x


Received <TapEvent Mutation, key="tap-1" (1 bytes
Value: x

Run Kill Close

○ From TAP to DCP



- TAP nearing end of life
- With 3.0 comes DCP (Database Change Protocol)
- Go SDK will have one of the first DCP implementations
- DCP only supported for internal replication at this time

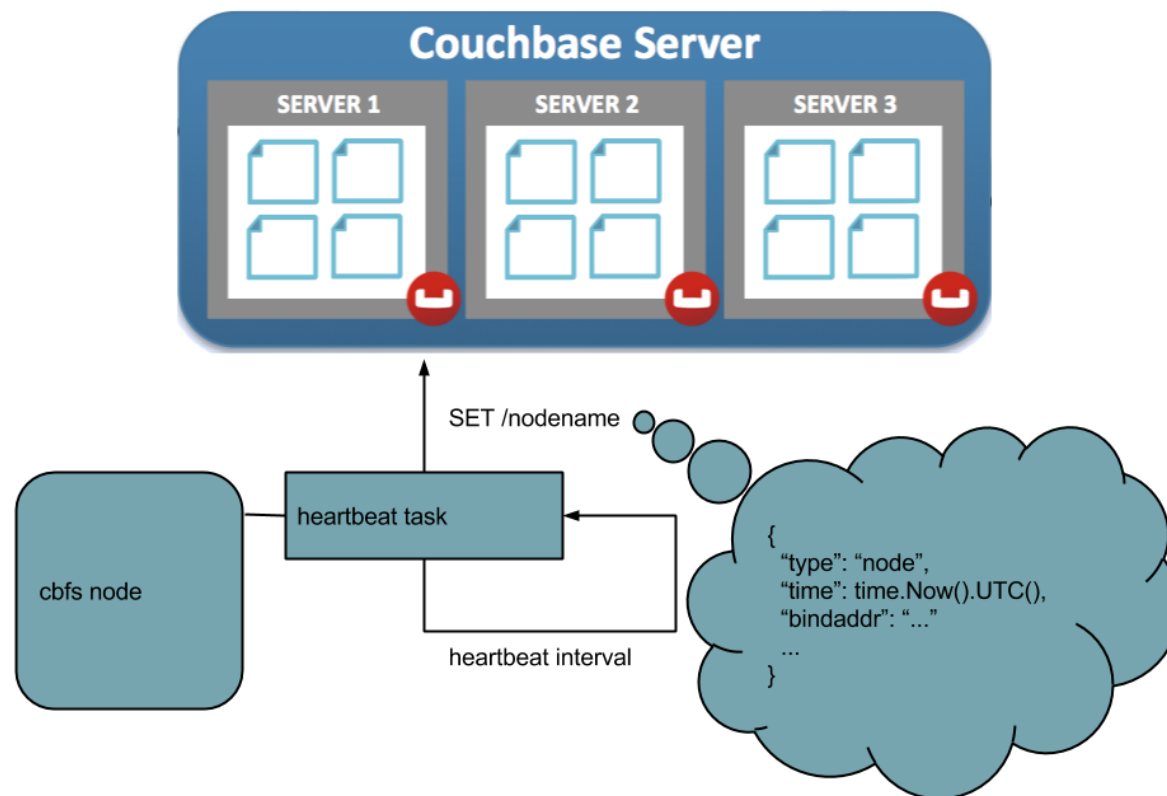


- Distributed file storage on top of Couchbase

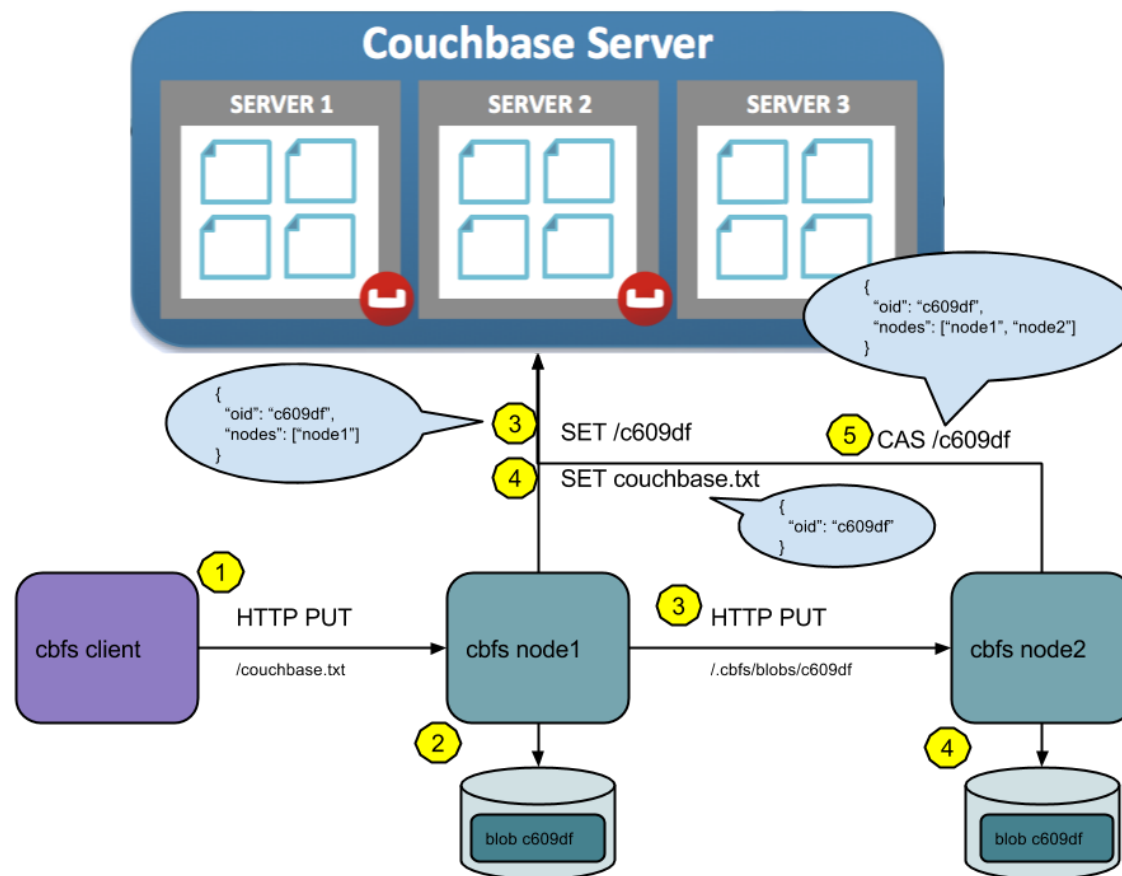


- Clients upload/download files via HTTP
- Nodes store file content locally in a content-addressable store (filename = content hash)
- File metadata is stored in Couchbase
- Nodes announce themselves/discover one another through Couchbase
- Nodes ensure a minimum replica count is maintained to safely store data

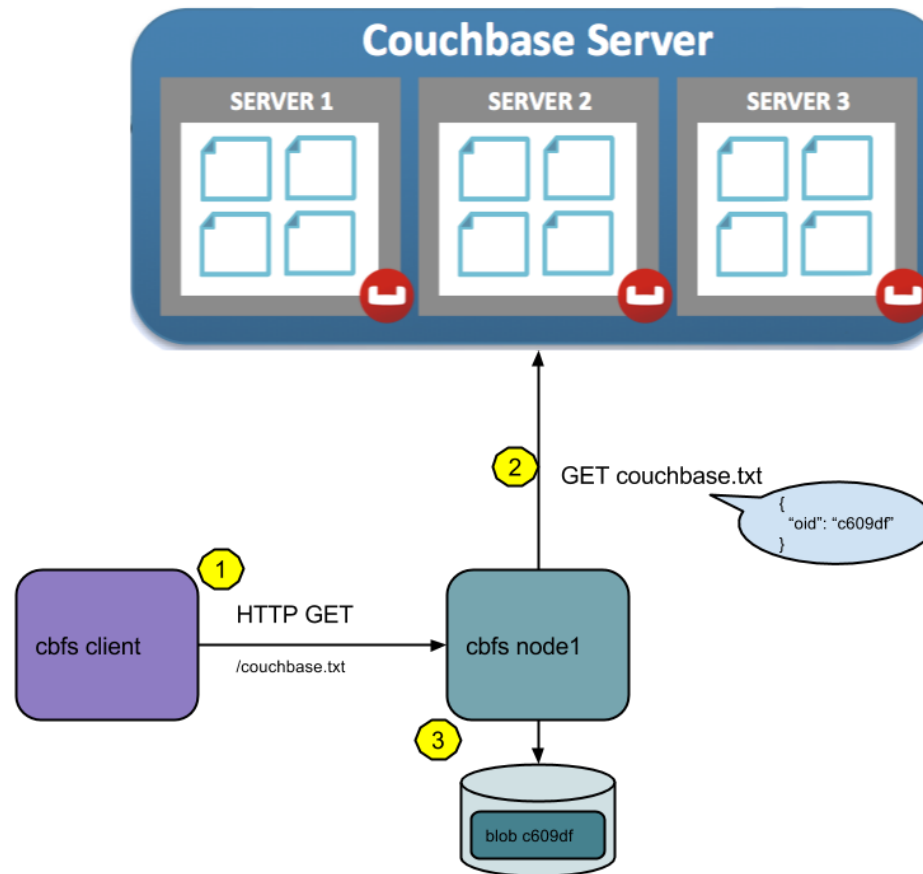
○ cbfs - Node Heartbeat



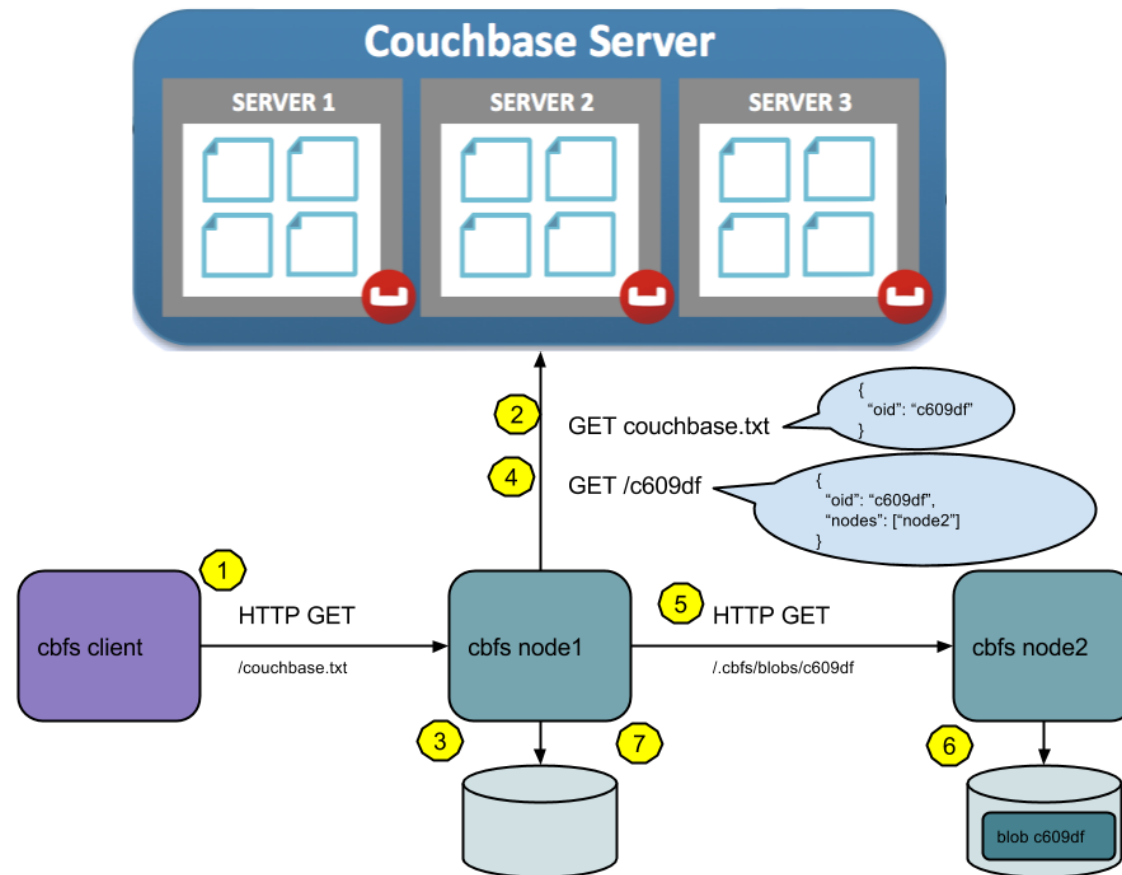
○ cbfs - Add Document



○ cbfs - Get Document (blob exist on node)




○ cbfs - Get Document (blob does NOT exist on node)





- Go - First class concurrency support with clean code
- Go - JSON mapping to custom structs
- Go - Out of the box support for HTTP/HTTPS
- Couchbase - Fast and scalable JSON storage
- Go + Couchbase = Powerful starting point for your app



○ Thank you



Marty Schoch

marty@couchbase.com (mailto:marty@couchbase.com)

<http://github.com/couchbaselabs/go-couchbase> (http://github.com/couchbaselabs/go-couchbase)

<https://github.com/couchbaselabs/cbugg> (https://github.com/couchbaselabs/cbugg)

https://github.com/couchbase/sync_gateway (https://github.com/couchbase/sync_gateway)

<https://github.com/couchbaselabs/cbfs> (https://github.com/couchbaselabs/cbfs)

[@mschoch](http://twitter.com/mschoch) (http://twitter.com/mschoch)